

```
In [131]: import pandas as pd                #importing pandas
          y=pd.read_csv('/home/placement/Desktop/college training/fiat500.csv') #reading csv using path
```

```
In [184]: y.groupby(['price']).head()        #grouping by price
```

Out[184]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price
0	1	lounge	51	882	25000	1	44.907242	8.611560	8900
1	2	pop	51	1186	32500	1	45.666359	12.241890	8800
2	3	sport	74	4658	142228	1	45.503300	11.417840	4200
3	4	lounge	51	2739	160000	1	40.633171	17.634609	6000
4	5	pop	73	3074	106880	1	41.903221	12.495650	5700
...	...	...	...	...	...	...	...	...	...
1515	1516	lounge	51	1917	124999	1	45.564491	10.115610	6999
1521	1522	lounge	51	3774	85000	1	44.294300	9.674440	4000
1523	1524	pop	51	2251	79800	1	45.512051	10.427010	6450
1531	1532	sport	73	4505	127000	1	45.528511	9.593230	4750
1534	1535	lounge	74	3835	112000	1	45.845692	8.666870	4600

657 rows × 9 columns

```
In [87]: h=y.rename(columns={'model':"model_name"})
h
```

```
# TO RENAME THE DATA
```

```
Out[87]:
```

	ID	model_name	engine_power	age_in_days	km	previous_owners	lat	lon	price
0	1	lounge	51	882	25000	1	44.907242	8.611560	8900
1	2	pop	51	1186	32500	1	45.666359	12.241890	8800
2	3	sport	74	4658	142228	1	45.503300	11.417840	4200
3	4	lounge	51	2739	160000	1	40.633171	17.634609	6000
4	5	pop	73	3074	106880	1	41.903221	12.495650	5700
...	...	...	...	...	...	...	...	...	...
1533	1534	sport	51	3712	115280	1	45.069679	7.704920	5200
1534	1535	lounge	74	3835	112000	1	45.845692	8.666870	4600
1535	1536	pop	51	2223	60457	1	45.481541	9.413480	7500
1536	1537	lounge	51	2557	80750	1	45.000702	7.682270	5990
1537	1538	pop	51	1766	54276	1	40.323410	17.568270	7900

1538 rows × 9 columns

In [88]: `y.describe()`

Out[88]:

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	price
<b>count</b>	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000
<b>mean</b>	769.500000	51.904421	1650.980494	53396.011704	1.123537	43.541361	11.563428	8576.003901
<b>std</b>	444.126671	3.988023	1289.522278	40046.830723	0.416423	2.133518	2.328190	1939.958641
<b>min</b>	1.000000	51.000000	366.000000	1232.000000	1.000000	36.855839	7.245400	2500.000000
<b>25%</b>	385.250000	51.000000	670.000000	20006.250000	1.000000	41.802990	9.505090	7122.500000
<b>50%</b>	769.500000	51.000000	1035.000000	39031.000000	1.000000	44.394096	11.869260	9000.000000
<b>75%</b>	1153.750000	51.000000	2616.000000	79667.750000	1.000000	45.467960	12.769040	10000.000000
<b>max</b>	1538.000000	77.000000	4658.000000	235000.000000	4.000000	46.795612	18.365520	11100.000000

In [185]: `y.head()`

Out[185]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price
<b>0</b>	1	lounge	51	882	25000	1	44.907242	8.611560	8900
<b>1</b>	2	pop	51	1186	32500	1	45.666359	12.241890	8800
<b>2</b>	3	sport	74	4658	142228	1	45.503300	11.417840	4200
<b>3</b>	4	lounge	51	2739	160000	1	40.633171	17.634609	6000
<b>4</b>	5	pop	73	3074	106880	1	41.903221	12.495650	5700

```
In [90]: x=y.drop("model",axis=1)           # to delete the data
x
```

Out[90]:

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	price
0	1	51	882	25000	1	44.907242	8.611560	8900
1	2	51	1186	32500	1	45.666359	12.241890	8800
2	3	74	4658	142228	1	45.503300	11.417840	4200
3	4	51	2739	160000	1	40.633171	17.634609	6000
4	5	73	3074	106880	1	41.903221	12.495650	5700
...	...	...	...	...	...	...	...	...
1533	1534	51	3712	115280	1	45.069679	7.704920	5200
1534	1535	74	3835	112000	1	45.845692	8.666870	4600
1535	1536	51	2223	60457	1	45.481541	9.413480	7500
1536	1537	51	2557	80750	1	45.000702	7.682270	5990
1537	1538	51	1766	54276	1	40.323410	17.568270	7900

1538 rows × 8 columns

```
In [91]: cor=x.corr()  
cor
```


*# to find the correlation*

Out[91]:

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	price
<b>ID</b>	1.000000	-0.034059	-0.060753	-0.006537	0.007803	-0.058207	0.058941	0.028516
<b>engine_power</b>	-0.034059	1.000000	0.319190	0.285495	-0.005030	0.005721	-0.005032	-0.277235
<b>age_in_days</b>	-0.060753	0.319190	1.000000	0.833890	0.075775	0.062982	-0.042667	-0.893328
<b>km</b>	-0.006537	0.285495	0.833890	1.000000	0.097539	0.035519	0.004839	-0.859373
<b>previous_owners</b>	0.007803	-0.005030	0.075775	0.097539	1.000000	0.001697	-0.026836	-0.076274
<b>lat</b>	-0.058207	0.005721	0.062982	0.035519	0.001697	1.000000	-0.766646	-0.011733
<b>lon</b>	0.058941	-0.005032	-0.042667	0.004839	-0.026836	-0.766646	1.000000	-0.003541
<b>price</b>	0.028516	-0.277235	-0.893328	-0.859373	-0.076274	-0.011733	-0.003541	1.000000

```
In [92]: import seaborn as sns
sns.heatmap(cor, vmax=1, vmin=-1, annot=True, linewidths=.5, cmap="bwr") # to print the heatmap
```

Out[92]: <Axes: >



```
In [93]: cor.groupby(["price", "km"]).count()      #correlation of the price and km
```

Out[93]:

		ID	engine_power	age_in_days	previous_owners	lat	lon
price	km						
-0.893328	0.833890	1	1	1	1	1	1
-0.859373	1.000000	1	1	1	1	1	1
-0.277235	0.285495	1	1	1	1	1	1
-0.076274	0.097539	1	1	1	1	1	1
-0.011733	0.035519	1	1	1	1	1	1
-0.003541	0.004839	1	1	1	1	1	1
0.028516	-0.006537	1	1	1	1	1	1
1.000000	-0.859373	1	1	1	1	1	1

```
In [94]: x['model']=y['model'].map({'lounge':1, 'pop':2, 'sport':4})    #renaming the terms
x
```

Out[94]:

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	price	model
0	1	51	882	25000	1	44.907242	8.611560	8900	1
1	2	51	1186	32500	1	45.666359	12.241890	8800	2
2	3	74	4658	142228	1	45.503300	11.417840	4200	4
3	4	51	2739	160000	1	40.633171	17.634609	6000	1
4	5	73	3074	106880	1	41.903221	12.495650	5700	2
...	...	...	...	...	...	...	...	...	...
1533	1534	51	3712	115280	1	45.069679	7.704920	5200	4
1534	1535	74	3835	112000	1	45.845692	8.666870	4600	1
1535	1536	51	2223	60457	1	45.481541	9.413480	7500	2
1536	1537	51	2557	80750	1	45.000702	7.682270	5990	1
1537	1538	51	1766	54276	1	40.323410	17.568270	7900	2

1538 rows × 9 columns



```
In [95]: y.groupby(["model", "ID"]).count() # groupby the data
```

```
Out[95]:
```

		engine_power	age_in_days	km	previous_owners	lat	lon	price
model	ID							
lounge	1	1	1	1	1	1	1	1
	4	1	1	1	1	1	1	1
	7	1	1	1	1	1	1	1
	8	1	1	1	1	1	1	1
	12	1	1	1	1	1	1	1
...	...	...	...	...	...	...	...	...
sport	1499	1	1	1	1	1	1	1
	1502	1	1	1	1	1	1	1
	1505	1	1	1	1	1	1	1
	1532	1	1	1	1	1	1	1
	1534	1	1	1	1	1	1	1

1538 rows × 7 columns

```
In [96]: a=x.sort_values("model")      # to sort the data
a
```

Out[96]:

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	price	model
0	1	51	882	25000	1	44.907242	8.61156	8900	1
952	953	51	2009	35552	1	42.562832	12.64220	8900	1
951	952	51	3684	78000	1	43.552910	10.31998	6500	1
950	951	51	762	36463	1	41.107880	14.20881	10900	1
949	950	51	790	32000	1	41.572239	13.33369	9490	1
...	...	...	...	...	...	...	...	...	...
1057	1058	51	3500	60000	2	45.438301	10.99170	6900	4
77	78	51	2739	77149	3	44.754890	8.03190	7800	4
172	173	51	4077	124000	1	41.152061	15.08309	4250	4
33	34	51	3927	140000	2	40.755932	14.69019	5200	4
775	776	51	2588	51000	1	45.536591	10.23204	7900	4

1538 rows × 9 columns

```
In [97]: y.groupby(['model']).count()
```

Out[97]:

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	price
model								
lounge	1094	1094	1094	1094	1094	1094	1094	1094
pop	358	358	358	358	358	358	358	358
sport	86	86	86	86	86	86	86	86

In [98]: `a.tail(5)`

Out[98]:

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	price	model
<b>1057</b>	1058	51	3500	60000	2	45.438301	10.99170	6900	4
<b>77</b>	78	51	2739	77149	3	44.754890	8.03190	7800	4
<b>172</b>	173	51	4077	124000	1	41.152061	15.08309	4250	4
<b>33</b>	34	51	3927	140000	2	40.755932	14.69019	5200	4
<b>775</b>	776	51	2588	51000	1	45.536591	10.23204	7900	4

In [99]: `a.head(5)`

Out[99]:

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	price	model
<b>0</b>	1	51	882	25000	1	44.907242	8.61156	8900	1
<b>952</b>	953	51	2009	35552	1	42.562832	12.64220	8900	1
<b>951</b>	952	51	3684	78000	1	43.552910	10.31998	6500	1
<b>950</b>	951	51	762	36463	1	41.107880	14.20881	10900	1
<b>949</b>	950	51	790	32000	1	41.572239	13.33369	9490	1

```
In [100]: data2=y
data2
```

Out[100]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price
0	1	lounge	51	882	25000	1	44.907242	8.611560	8900
1	2	pop	51	1186	32500	1	45.666359	12.241890	8800
2	3	sport	74	4658	142228	1	45.503300	11.417840	4200
3	4	lounge	51	2739	160000	1	40.633171	17.634609	6000
4	5	pop	73	3074	106880	1	41.903221	12.495650	5700
...	...	...	...	...	...	...	...	...	...
1533	1534	sport	51	3712	115280	1	45.069679	7.704920	5200
1534	1535	lounge	74	3835	112000	1	45.845692	8.666870	4600
1535	1536	pop	51	2223	60457	1	45.481541	9.413480	7500
1536	1537	lounge	51	2557	80750	1	45.000702	7.682270	5990
1537	1538	pop	51	1766	54276	1	40.323410	17.568270	7900

1538 rows × 9 columns

```
In [101]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test =train_test_split(x,y,test_size=0.33,random_state=42) #to import the sklearn
```

```
In [166]: from sklearn.linear_model import LinearRegression
reg= LinearRegression () #creatuig objects of linear regression
reg.fit(x_train,y_train) #training and fitting lr objects training data
```

Out[166]:

```
▼ LinearRegression
LinearRegression()
```

```
In [167]: pred=reg.predict(x_test)
pred
```

```
8426.37409382, 5447.47569851, 9914.20077691, 4687.39013431,
7885.32100747, 5431.00822998, 9911.86294348, 10390.16991322,
9680.84745901, 8844.57815539, 7764.08471024, 4257.54640953,
9882.76503303, 10341.35258769, 5736.4484335 , 10179.87154436,
9501.423448 , 7997.3181334 , 5532.33458288, 9894.57834738,
10437.97459358, 6381.35845844, 9591.23555726, 9574.27908517,
10322.30715736, 9501.22785499, 9789.955758 , 9593.26549752,
6775.82788536, 7915.34831306, 10389.98590521, 10351.58343315,
7381.32686464, 9966.53983093, 10430.87188433, 10554.43156462,
10285.85574963, 10035.88086558, 9526.63034431, 7742.78157141,
9297.64938364, 10051.42272678, 10004.81256571, 9985.84167026,
9374.6573594 , 9561.57499854, 9754.94184269, 9819.85893758,
8780.31447831, 6255.99008069, 6281.53627686, 8190.88781577,
8588.91394592, 6566.97963218, 6850.70237466, 5511.29438169,
8119.97866315, 9847.74830838, 7775.93862032, 9875.05509733,
10121.29366536, 5791.92464084, 9835.42728501, 10043.91426822,
8027.28015259, 4527.22080416, 10609.02444098, 3808.29240951,
9952.37340054, 10511.20945172, 5746.34019592, 5486.40214756,
10395.91036208, 6788.47519216, 8953.20120295, 10442.24187982,
9455.6934072 . 9976.26574762. 8528.35753837. 7960.77147517.
```

```
In [102]: x_test.head(10)
```

```
Out[102]:
```

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	price	model
<b>481</b>	482	51	3197	120000	2	40.174702	18.167629	7900	2
<b>76</b>	77	62	2101	103000	1	45.797859	8.644440	7900	2
<b>1502</b>	1503	51	670	32473	1	41.107880	14.208810	9400	1
<b>669</b>	670	51	913	29000	1	45.778591	8.946250	8500	1
<b>1409</b>	1410	51	762	18800	1	45.538689	9.928310	9700	1
<b>1414</b>	1415	51	762	39751	1	41.903221	12.495650	9900	1
<b>1089</b>	1090	51	882	33160	1	45.778999	12.997090	9900	1
<b>1507</b>	1508	51	701	17324	1	45.556549	9.534470	9950	1
<b>970</b>	971	51	701	29000	1	36.855839	14.760470	10700	1
<b>1198</b>	1199	51	1155	38000	1	41.239281	13.933020	8999	1

```
In [103]: a=y.drop(["ID","lat","lon"],axis=1)          #dropping the data  
a
```

Out[103]:

	model	engine_power	age_in_days	km	previous_owners	price
0	lounge	51	882	25000	1	8900
1	pop	51	1186	32500	1	8800
2	sport	74	4658	142228	1	4200
3	lounge	51	2739	160000	1	6000
4	pop	73	3074	106880	1	5700
...	...	...	...	...	...	...
1533	sport	51	3712	115280	1	5200
1534	lounge	74	3835	112000	1	4600
1535	pop	51	2223	60457	1	7500
1536	lounge	51	2557	80750	1	5990
1537	pop	51	1766	54276	1	7900

1538 rows × 6 columns

```
In [104]: a=pd.get_dummies(a)           # to print the dummies
a
```

Out[104]:

	engine_power	age_in_days	km	previous_owners	price	model_lounge	model_pop	model_sport
0	51	882	25000	1	8900	1	0	0
1	51	1186	32500	1	8800	0	1	0
2	74	4658	142228	1	4200	0	0	1
3	51	2739	160000	1	6000	1	0	0
4	73	3074	106880	1	5700	0	1	0
...	...	...	...	...	...	...	...	...
1533	51	3712	115280	1	5200	0	0	1
1534	74	3835	112000	1	4600	1	0	0
1535	51	2223	60457	1	7500	0	1	0
1536	51	2557	80750	1	5990	1	0	0
1537	51	1766	54276	1	7900	0	1	0

1538 rows × 8 columns



```
In [105]: b=pd.get_dummies(a)
b
```

```
Out[105]:
```

	engine_power	age_in_days	km	previous_owners	price	model_lounge	model_pop	model_sport
0	51	882	25000	1	8900	1	0	0
1	51	1186	32500	1	8800	0	1	0
2	74	4658	142228	1	4200	0	0	1
3	51	2739	160000	1	6000	1	0	0
4	73	3074	106880	1	5700	0	1	0
...	...	...	...	...	...	...	...	...
1533	51	3712	115280	1	5200	0	0	1
1534	74	3835	112000	1	4600	1	0	0
1535	51	2223	60457	1	7500	0	1	0
1536	51	2557	80750	1	5990	1	0	0
1537	51	1766	54276	1	7900	0	1	0

1538 rows × 8 columns

```
In [106]: b.shape
```

*# to find the shape in row and columns*

```
Out[106]: (1538, 8)
```

```
In [107]: y=b["price"]
x=b.drop('price',axis=1)
```

*# to delete the rows*

In [108]:

y

```
Out[108]: 0      8900
          1      8800
          2      4200
          3      6000
          4      5700
          ...
          1533    5200
          1534    4600
          1535    7500
          1536    5990
          1537    7900
```

Name: price, Length: 1538, dtype: int64

```
In [109]: from sklearn.model_selection import train_test_split
          x_train,x_test,y_train,y_test =train_test_split(x,y,test_size=0.33,random_state=42) #to import the sklearn
```

In [110]: x\_test.head(10)

Out[110]:

	engine_power	age_in_days	km	previous_owners	model_lounge	model_pop	model_sport
481	51	3197	120000	2	0	1	0
76	62	2101	103000	1	0	1	0
1502	51	670	32473	1	1	0	0
669	51	913	29000	1	1	0	0
1409	51	762	18800	1	1	0	0
1414	51	762	39751	1	1	0	0
1089	51	882	33160	1	1	0	0
1507	51	701	17324	1	1	0	0
970	51	701	29000	1	1	0	0
1198	51	1155	38000	1	1	0	0

```
In [111]: y_test.head(10)
```

```
Out[111]: 481      7900  
          76      7900  
          1502    9400  
          669    8500  
          1409    9700  
          1414    9900  
          1089    9900  
          1507    9950  
          970   10700  
          1198    8999  
          Name: price, dtype: int64
```

```
In [112]: from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import Ridge
```

```
alpha = [1e-15, 1e-10, 1e-8, 1e-4, 1e-3, 1e-2, 1, 5, 10, 20, 30]
```

```
ridge = Ridge()
```

```
parameters = {'alpha': alpha}
```

```
ridge_regressor = GridSearchCV(ridge, parameters)
```

```
ridge_regressor.fit(x_train, y_train)
```

```
/home/placement/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_ridge.py:216: LinAlgWarning:
Ill-conditioned matrix (rcond=5.56109e-26): result may not be accurate.
```

```
    return linalg.solve(A, Xy, assume_a="pos", overwrite_a=True).T
```

```
/home/placement/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_ridge.py:216: LinAlgWarning:
Ill-conditioned matrix (rcond=7.70876e-26): result may not be accurate.
```

```
    return linalg.solve(A, Xy, assume_a="pos", overwrite_a=True).T
```

```
/home/placement/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_ridge.py:216: LinAlgWarning:
Ill-conditioned matrix (rcond=6.91585e-23): result may not be accurate.
```

```
    return linalg.solve(A, Xy, assume_a="pos", overwrite_a=True).T
```

```
/home/placement/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_ridge.py:216: LinAlgWarning:
Ill-conditioned matrix (rcond=7.08003e-23): result may not be accurate.
```

```
    return linalg.solve(A, Xy, assume_a="pos", overwrite_a=True).T
```

```
/home/placement/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_ridge.py:216: LinAlgWarning:
Ill-conditioned matrix (rcond=7.01022e-23): result may not be accurate.
```

```
    return linalg.solve(A, Xy, assume_a="pos", overwrite_a=True).T
```

```
/home/placement/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_ridge.py:216: LinAlgWarning:
Ill-conditioned matrix (rcond=7.57959e-23): result may not be accurate.
```

```
    return linalg.solve(A, Xy, assume_a="pos", overwrite_a=True).T
```

```
/home/placement/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_ridge.py:216: LinAlgWarning:
Ill-conditioned matrix (rcond=7.24161e-23): result may not be accurate.
```

```
    return linalg.solve(A, Xy, assume_a="pos", overwrite_a=True).T
```

```
/home/placement/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_ridge.py:216: LinAlgWarning:
Ill-conditioned matrix (rcond=6.92759e-21): result may not be accurate.
```

```
    return linalg.solve(A, Xy, assume_a="pos", overwrite_a=True).T
```

```
/home/placement/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_ridge.py:216: LinAlgWarning:
Ill-conditioned matrix (rcond=7.09091e-21): result may not be accurate.
```

```
    return linalg.solve(A, Xy, assume_a="pos", overwrite_a=True).T
```


```
/home/placement/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_ridge.py:216: LinAlgWarning:
```

```

Ill-conditioned matrix (rcond=7.02112e-21): result may not be accurate.
    return linalg.solve(A, Xy, assume_a="pos", overwrite_a=True).T
/home/placement/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_ridge.py:216: LinAlgWarning:
Ill-conditioned matrix (rcond=7.57414e-21): result may not be accurate.
    return linalg.solve(A, Xy, assume_a="pos", overwrite_a=True).T
/home/placement/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_ridge.py:216: LinAlgWarning:
Ill-conditioned matrix (rcond=7.23284e-21): result may not be accurate.
    return linalg.solve(A, Xy, assume_a="pos", overwrite_a=True).T
/home/placement/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_ridge.py:216: LinAlgWarning:
Ill-conditioned matrix (rcond=6.9277e-17): result may not be accurate.
    return linalg.solve(A, Xy, assume_a="pos", overwrite_a=True).T
/home/placement/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_ridge.py:216: LinAlgWarning:
Ill-conditioned matrix (rcond=7.09099e-17): result may not be accurate.
    return linalg.solve(A, Xy, assume_a="pos", overwrite_a=True).T
/home/placement/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_ridge.py:216: LinAlgWarning:
Ill-conditioned matrix (rcond=7.02123e-17): result may not be accurate.
    return linalg.solve(A, Xy, assume_a="pos", overwrite_a=True).T
/home/placement/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_ridge.py:216: LinAlgWarning:
Ill-conditioned matrix (rcond=7.57407e-17): result may not be accurate.
    return linalg.solve(A, Xy, assume_a="pos", overwrite_a=True).T
/home/placement/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_ridge.py:216: LinAlgWarning:
Ill-conditioned matrix (rcond=7.23274e-17): result may not be accurate.
    return linalg.solve(A, Xy, assume_a="pos", overwrite_a=True).T

```

Out[112]:



In [113]: `ridge_regressor.best_params_`

Out[113]: `{'alpha': 30}`

```
In [116]: from sklearn.metrics import mean_squared_error           #calculating the mse
```

```
In [114]: ridge=Ridge(alpha=30)
          ridge.fit(x_train,y_train)
          pred_ridge=ridge.predict(x_test)
```

```
In [119]: ridge
```

```
Out[119]: 

▼      Ridge
  Ridge(alpha=30)


```

```
In [117]: Ridge_Error=mean_squared_error(pred_ridge,y_test)
          Ridge_Error                                           #calculation of ridg error
```

```
Out[117]: 579521.7970897449
```

```
In [182]: from sklearn.metrics import r2_score
          r2_score(y_test,pred_ridge)
```

```
Out[182]: 0.8421969385523054
```

**# NEXT**

```
In [129]: y.groupby(['price']).count()
```

Out[129]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon
price								
2500	1	1	1	1	1	1	1	1
2900	1	1	1	1	1	1	1	1
3390	1	1	1	1	1	1	1	1
3500	1	1	1	1	1	1	1	1
3600	1	1	1	1	1	1	1	1
...	...	...	...	...	...	...	...	...
10990	9	9	9	9	9	9	9	9
10999	5	5	5	5	5	5	5	5
11000	13	13	13	13	13	13	13	13
11090	2	2	2	2	2	2	2	2
11100	1	1	1	1	1	1	1	1

222 rows × 8 columns

```
In [142]: b=y.loc[(y.model=="lounge")]
```

In [143]:

b

Out[143]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price
0	1	lounge	51	882	25000	1	44.907242	8.611560	8900
3	4	lounge	51	2739	160000	1	40.633171	17.634609	6000
6	7	lounge	51	731	11600	1	44.907242	8.611560	10750
7	8	lounge	51	1521	49076	1	41.903221	12.495650	9190
11	12	lounge	51	366	17500	1	45.069679	7.704920	10990
...	...	...	...	...	...	...	...	...	...
1528	1529	lounge	51	2861	126000	1	43.841980	10.515310	5500
1529	1530	lounge	51	731	22551	1	38.122070	13.361120	9900
1530	1531	lounge	51	670	29000	1	45.764648	8.994500	10800
1534	1535	lounge	74	3835	112000	1	45.845692	8.666870	4600
1536	1537	lounge	51	2557	80750	1	45.000702	7.682270	5990

1094 rows × 9 columns



```
In [145]: c=l.drop("model",axis=1)           # to delete the data
          c
```

Out[145]:

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	price
0	1	51	882	25000	1	44.907242	8.611560	8900
3	4	51	2739	160000	1	40.633171	17.634609	6000
6	7	51	731	11600	1	44.907242	8.611560	10750
7	8	51	1521	49076	1	41.903221	12.495650	9190
11	12	51	366	17500	1	45.069679	7.704920	10990
...	...	...	...	...	...	...	...	...
1528	1529	51	2861	126000	1	43.841980	10.515310	5500
1529	1530	51	731	22551	1	38.122070	13.361120	9900
1530	1531	51	670	29000	1	45.764648	8.994500	10800
1534	1535	74	3835	112000	1	45.845692	8.666870	4600
1536	1537	51	2557	80750	1	45.000702	7.682270	5990

1094 rows × 8 columns

```
In [149]: cory=c.corr()  
cory
```


*# to find the correlation*

Out[149]:

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	price
ID	1.000000	-0.007859	-0.055600	0.008471	0.000486	-0.067365	0.085564	0.004406
engine_power	-0.007859	1.000000	0.300144	0.281481	-0.027861	-0.027428	0.025088	-0.260315
age_in_days	-0.055600	0.300144	1.000000	0.837210	0.038825	0.059653	-0.060897	-0.892417
km	0.008471	0.281481	0.837210	1.000000	0.050958	0.019335	0.007749	-0.859062
previous_owners	0.000486	-0.027861	0.038825	0.050958	1.000000	0.014027	-0.052962	-0.039448
lat	-0.067365	-0.027428	0.059653	0.019335	0.014027	1.000000	-0.769521	0.022973
lon	0.085564	0.025088	-0.060897	0.007749	-0.052962	-0.769521	1.000000	-0.014498
price	0.004406	-0.260315	-0.892417	-0.859062	-0.039448	0.022973	-0.014498	1.000000

```
In [164]: import seaborn as sns
import matplotlib.pyplot as plt
sns.heatmap(cory, vmax=1, vmin=-1, annot=True, linewidths=.5, cmap="bwr") # to print the heatmap
```

Out[164]: <Axes: >



```
In [181]: re=pd.DataFrame(columns=['actual','predicted'])           #to compare the values
re['actual']=y_test
re['predicted']=pred
re=re.reset_index()
re['Id']=re.index          # to arrange the indexing

re.head(5)
```

Out[181]:

	index	actual	predicted	Id
0	481	7900	5867.650338	0
1	76	7900	7133.701423	1
2	1502	9400	9866.357762	2
3	669	8500	9723.288745	3
4	1409	9700	10039.591012	4

```
In [188]: import seaborn as sns
import matplotlib.pyplot as plt

sns.lineplot(x='Id',y='actual',data=re.head(100))
sns.lineplot(x='Id',y='predicted',data=re.head(100))    #to plot the line graph
plt.plot()
```

Out[188]: []

